Tommaso **Peresson**

Data Platform Engineer

(+39) 340-2905127 | tompere96@gmail.com | t0mpere | tommaso-peresson

Summary

Data Platform Engineer with 6 years building PB-scale batch and streaming infrastructure (Spark, Kafka, Flink, Pinot, Iceberg, Delta Lake). Now at Depop, cutting platform cost and building AI developer tooling; previously technical lead on Bumble's data platform. Apache Pinot contributor. MSc, Politecnico di Milano.

Experience

Depop — *London*

Data Platform Engineer — *Feb 2026 – Present*

- Led the cost optimisation of the company's largest Spark ingestion pipeline (Databricks, Delta Lake), cutting its cost by **~52%** without degrading downstream query performance (**blog post**).
- Built an AI-powered self-service toolkit for A/B-test feature flags (Claude Code, Backstage, GitHub Actions, Python), cutting flag creation time by **~97%**; adopted for 50% of new flags within weeks (**blog post**).

Bumble Inc. — *London*

Data Platform Engineer — *Apr 2022 – Jan 2026*

- Designed a cloud lakehouse architecture (Iceberg, Polaris, Kafka, Flink, Spark on GCP) to migrate a 200+ node on-prem Hadoop/YARN cluster.
- Technical leadership for the data platform team; mentored Data and Analytics Engineers and led knowledge sharing across the data org.
- Built in-house platform tooling: DDL schema generators, schema registry, client SDKs, orchestrators and UIs (Java, Python, TS, PHP).
- Built near-real-time apps for anomaly detection, A/B testing and data discovery (Pinot-based, **10x+ faster** pages).
- Maintained hybrid PB-scale data infrastructure on the on-call rotation (Kafka, Hadoop, Spark, Snowflake, ClickHouse, Pinot).

HousingAnywhere — *Rotterdam*

Data Engineer — *Mar 2021 – Mar 2022*

- Led the restructuring of the Snowflake data warehouse schema.
- Designed and implemented the data warehouse permission model with growth and scalability in mind.
- Designed and built the warehouse **anonymisation layer** using dbt (**published in Towards Data Science**).
- Designed and implemented ETL and ELT pipelines with Airflow.

HousingAnywhere — *Rotterdam*

Data Engineering Intern — *Aug 2020 – Feb 2021*

- Data engineering and data science; MSc thesis on a state-of-the-art recommender system for a two-sided marketplace.

Skills

Languages	Java, Scala, Python, SQL, TypeScript, PHP, Bash
Streaming & Processing	Kafka, Flink, Spark, Hadoop/YARN
Storage & Query	Iceberg, Delta Lake, Pinot, ClickHouse, Snowflake, Hive
Cloud & Ops	Databricks, GCP, Kubernetes, Docker, Airflow, dbt, GitHub Actions, Claude Code
Spoken	Italian (native), English (C1)

Open Source

Apache Pinot

Contributor – 3 merged PRs

- Star-tree index compression for functionColumnPairs (#11744, 2023).
- Configuration through environment variables (#12307, 2024).
- MergeRollupTask output-dir config fix (#17015, 2025).

Education

Politecnico di Milano — *Milan*

MSc & BSc in Computer Science Engineering — *2015 – 2021*

Google Cloud

Professional Data Engineer Certification — *Jan 2022*